


# Why recommendation change is not automatically unfair

FairEval keeps the task and preference evidence fixed, changes one registered context factor, then asks what the ranking change does to held-out utility.



**Core principle:  $\Delta$  ranking  $\neq$  unfairness. Fairness claims require matched interventions + preference-conditioned consequence evidence + robustness.**